

Low-Power Neural Network Accelerators: Advancements in Custom Floating-Point Techniques

Yarib Nevarez

Universität Bremen

May 22, 2024



Introduction

- Advancements in AI/ML for IoT and TinyML

Introduction

- Advancements in AI/ML for IoT and TinyML
- Need for energy efficiency

Introduction

- Advancements in AI/ML for IoT and TinyML
- Need for energy efficiency
- Hybrid precision approach

Introduction

- Advancements in AI/ML for IoT and TinyML
- Need for energy efficiency
- Hybrid precision approach
- Design challenges

Introduction

- Advancements in AI/ML for IoT and TinyML
- Need for energy efficiency
- Hybrid precision approach
- Design challenges
- Approximate computing and quantization

Introduction

- Advancements in AI/ML for IoT and TinyML
- Need for energy efficiency
- Hybrid precision approach
- Design challenges
- Approximate computing and quantization
- Quality, interoperability, and compatibility

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation
 - Energy-efficient design strategies and custom FP arithmetic units

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation
 - Energy-efficient design strategies and custom FP arithmetic units
 - Precision and quantization impact analysis

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation
 - Energy-efficient design strategies and custom FP arithmetic units
 - Precision and quantization impact analysis
 - Hardware performance evaluation and bench marking

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation
 - Energy-efficient design strategies and custom FP arithmetic units
 - Precision and quantization impact analysis
 - Hardware performance evaluation and bench marking
 - Practical application

Goal and Objectives

- Goal: To develop advanced methodologies for energy-efficient neural network accelerators with custom floating-point (FP) computation in low-power and resource constrained devices.
- Objectives:
 - Optimal custom FP number representation
 - Energy-efficient design strategies and custom FP arithmetic units
 - Precision and quantization impact analysis
 - Hardware performance evaluation and bench marking
 - Practical application
 - Future research

Outline

1 State-of-the-Art

2 Hybrid 8-bit Floating-Point and 4-bit Logarithmic Computation

1 State-of-the-Art

2 Hybrid 8-bit Floating-Point and 4-bit Logarithmic Computation

High-Performance FPGA-Based CNN Accelerator With Block-Floating-Point Arithmetic

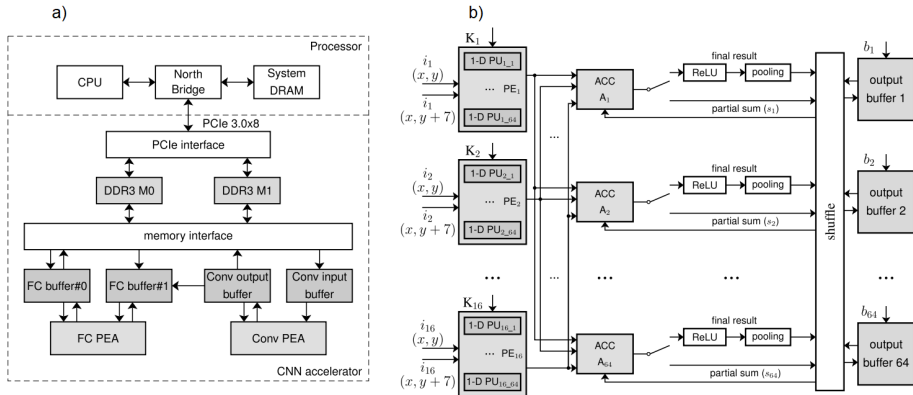
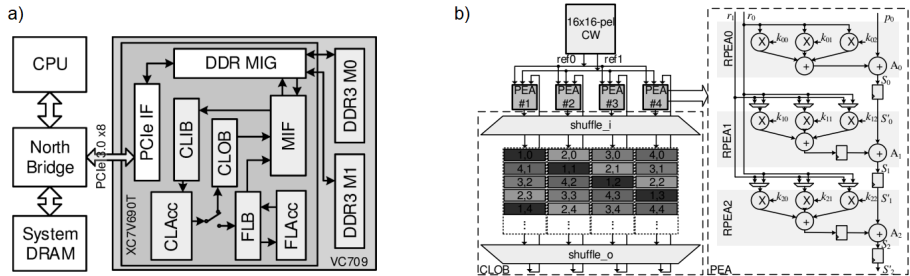


Figure: (a) System architecture. (b) Processing element array.

A 200MHZ 202.4GFLOPS@10.8W VGG16 Accelerator in Xilinx VX690T



Low-precision Floating-point Arithmetic for High-performance FPGA-based CNN Acceleration

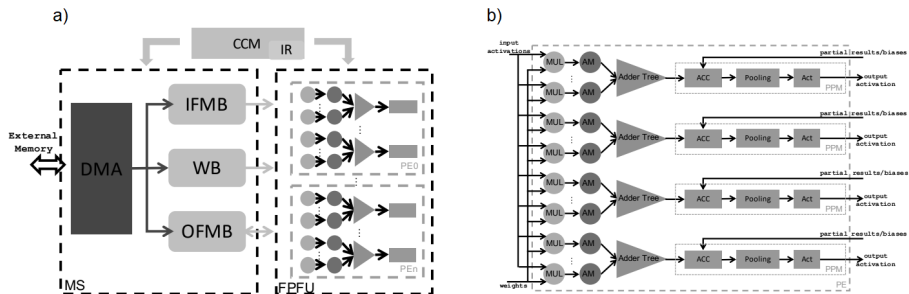


Figure: (a) System architecture. (b) Processing element.

CNN Hardware Acceleration on a Low-Power and Low-Cost APSoC

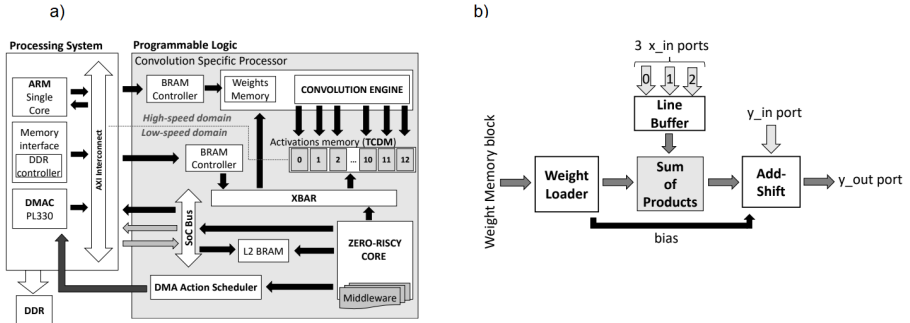


Figure: (a) System architecture. (b) Convolution engine.

1 State-of-the-Art

2 Hybrid 8-bit Floating-Point and 4-bit Logarithmic Computation ▶

Spike-by-Spike Neural Network




Figure: Spike-by-Spike (SbS) neural network architecture for handwritten digit classification task.

Spike-by-Spike Neural Network




Figure: Spike-by-Spike (SbS) neural network architecture for handwritten digit classification task.




Figure: Performance classification of SbS NN versus equivalent CNN.

Spike-by-Spike Layer Update




Figure: SbS inference population (IP) as independent computational entities.

$$h_{\mu}^{new}(i) = \frac{1}{1 + \epsilon} \left(h_{\mu}(i) + \epsilon \frac{h_{\mu}(i)W(s_t|i)}{\sum_j h_{\mu}(j)W(s_t|j)} \right)$$

HW/SW Co-Design Framework




Figure: System-level overview of the embedded software architecture.




Figure: System-level hardware architecture with scalable number of heterogeneous processing units (PU).

Processing Unit

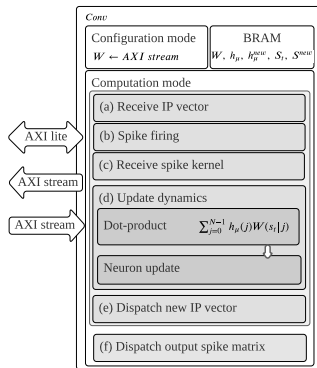


Figure: Conv processing unit.

Processing Unit

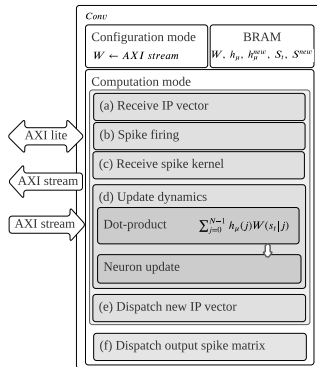


Figure: Conv processing unit.

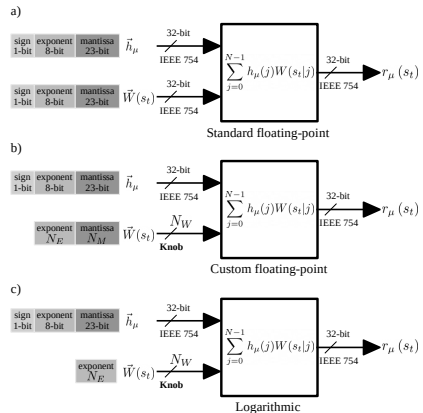


Figure: Dot-product hardware module.

Hybrid Dot-Product Approximation

$$r_{\mu}(s_t) = \sum_{j=0}^{N-1} h_{\mu}(j) W(s_t[j]) \quad (1)$$

$$E_{\min} = \log_2(\min_{\forall i}(W(i))) \quad (2)$$

$$N_E = \lceil \log_2(|E_{\min}|) \rceil \quad (3)$$

$$N_W = N_E + N_M \quad (4)$$

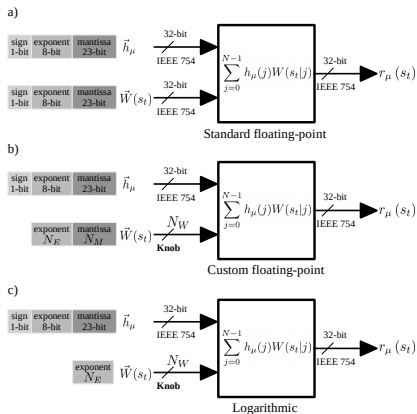


Figure: Dot-product hardware module.

Dot-Product with Standard Floating-Point (IEEE 754)

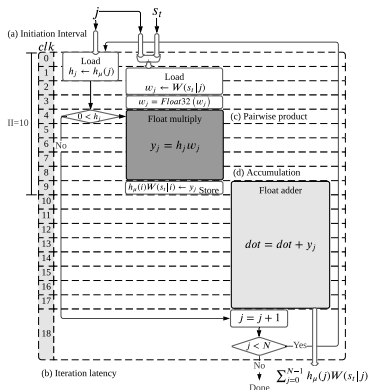


Figure: Dot-product hardware module with standard floating-point computation.

$$L_{f32} = 10N + 9$$

Dot-Product with Hybrid Custom Floating-Point Approximation




Figure: Dot-product hardware module with hybrid custom floating-point approximation.

$$L_{\text{custom}} = 2N + 11$$

Dot-product with Hybrid Logarithmic Approximation

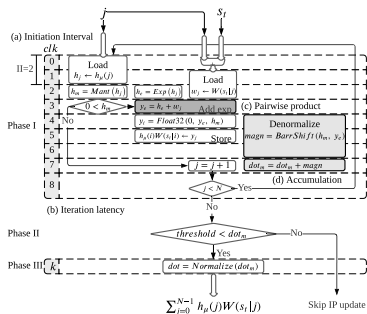


Figure: Dot-product hardware module with hybrid logarithmic approximation.

$$L_{\text{custom}} = 2N + 7$$

Acceleration with Standard Floating-Point




Figure: System overview of the top-level architecture with 8 processing units.

Acceleration with Standard Floating-Point




Figure: System overview of the top-level architecture with 8 processing units.




Figure: Computation on embedded CPU.

Acceleration with Standard Floating-Point




Figure: System overview of the top-level architecture with 8 processing units.




Figure: Performance of processing units with standard floating-point.




Figure: Computation on embedded CPU.

Acceleration with Standard Floating-Point




Figure: System overview of the top-level architecture with 8 processing units.




Figure: Computation on embedded CPU.




Figure: Performance of processing units with standard floating-point.




Figure: Performance bottleneck of cyclic computation on processing units with standard floating-point.

Acceleration with Custom Floating-Point

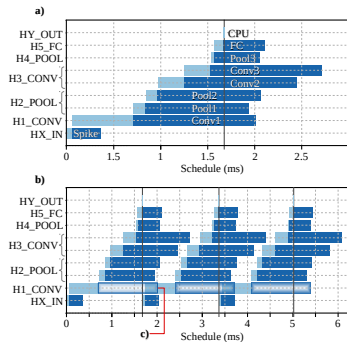


Figure: Performance on processing units with hybrid 8-bit floating-point.

Acceleration with Custom Floating-Point




Figure: Performance on processing units with hybrid 8-bit floating-point.




Figure: Performance of processing units with hybrid 4-bit logarithmic approximation.

Noise tolerance




Figure: Noise tolerance with standard floating-point.

Noise tolerance

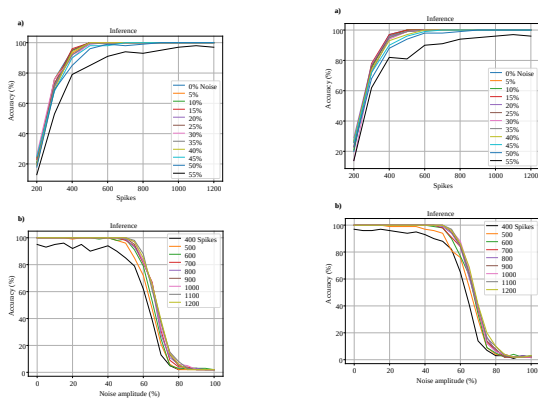


Figure: Noise tolerance with standard floating-point.

Figure: Noise tolerance with hybrid 8-bit floating-point approximation.

Noise tolerance




Figure: Noise tolerance with standard floating-point.




Figure: Noise tolerance with hybrid 8-bit floating-point approximation.




Figure: Noise tolerance with hybrid 4-bit logarithmic approximation.

Accelerator Implementations

Table: Accelerator implementations.

Platform implementation	Power (W)	Clk (MHz)	Latency (ms)	Acceleration	Accuracy (%)
Standard floating-point	2.420	200	3.18	10.7x	98.98
Hybrid floating-point 8-bit	2.369	200	1.67	20.5x	98.97
Hybrid Logarithmic 4-bit	2.324	200	1.67	20.5x	98.84