


Multi-Phase Model Optimization


Training with iterative early stop with Adam resets



Loss function = MSE

Early stop = 10 epochs
Mini batch = 512 samples

Quantization-aware training for hybrid floating-point 6-bit



Total time = 53 minutes

CPU = AMD Ryzen 5600H

GPU = NVIDIA GeForce RTX 3050