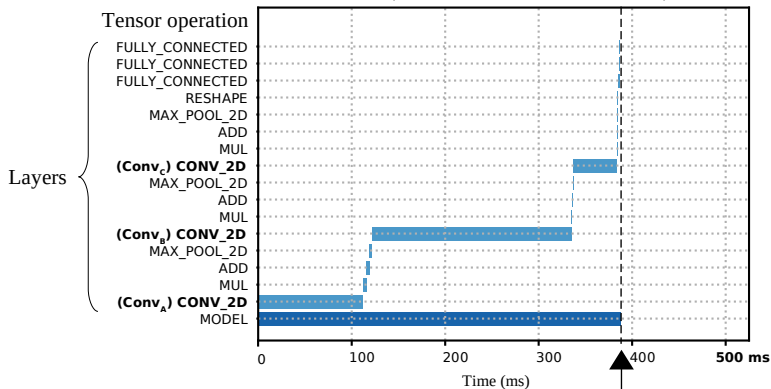**Runtime Profiling: Latencies and Scheduling on Floating-Point and Hybrid Float 6 Accelerators**

Tensor Op	CPU (ms)	TP FP32 (ms)	Gain	TP HF6 (ms)	Gain
$Conv_A$					
$Conv_B$					
$Conv_C$					

CPU (ARM Cortex-A9 @ 666 MHz)

Total = 387 ms (2.6 inference per second)